

# R Notebook

## Argus Media - Conding Test

Ivan Berlim Gonçalves

---

The purpose of this exercise is to design and implement an entire data preparation pipeline in R. We would like you to implement a robust, extensible and generic framework for data preparation.

---

```
# Verify if the package is already installed, if not, install package
packages <- c("gamlss", "gamlss.add", "gamlss.dist", "roll", "dplyr", "tseries", "ggpubr", "magrittr", "labelVector", "corrplot")
install.packages(setdiff(packages, rownames(installed.packages())))

# Loading libraries
library(gamlss)
library(gamlss.add)
library(gamlss.dist)
# library(DT)
library(roll)
library(dplyr)
# library(stats)
library(ggpubr)
library(tseries)
# library(ggplot2)
# library(dataspice)
library(Hmisc)
library(magrittr)
library(labelVector)
library(corrplot)
```

### Packages and Libraries

- 1) Take as raw inputs to the data preparation process, the oil data from the gamlss package.

```
# Raw data preparation
data_prep <- function(rawdata) {
  oil_data <- rawdata
  return(oil_data)
}
```

```

oil_data <- data_prep(gamlss.data::oil)

# Printing data information
paste0("The data set has ", nrow(oil_data), " observations and ", ncol(oil_data), " variables. ")

## [1] "The data set has 1000 observations and 25 variables. "

cat('\n')

oil_col_names <- colnames(oil_data)
cat('Columns in the data frame: \n\n')

## Columns in the data frame:

for (i in (1:ncol(oil))) {print(paste0("Column number: ",i, ". Variable name: ", oil_col_names[i]))}

## [1] "Column number: 1. Variable name: OILPRICE"
## [1] "Column number: 2. Variable name: CL2_log"
## [1] "Column number: 3. Variable name: CL3_log"
## [1] "Column number: 4. Variable name: CL4_log"
## [1] "Column number: 5. Variable name: CL5_log"
## [1] "Column number: 6. Variable name: CL6_log"
## [1] "Column number: 7. Variable name: CL7_log"
## [1] "Column number: 8. Variable name: CL8_log"
## [1] "Column number: 9. Variable name: CL9_log"
## [1] "Column number: 10. Variable name: CL10_log"
## [1] "Column number: 11. Variable name: CL11_log"
## [1] "Column number: 12. Variable name: CL12_log"
## [1] "Column number: 13. Variable name: CL13_log"
## [1] "Column number: 14. Variable name: CL14_log"
## [1] "Column number: 15. Variable name: CL15_log"
## [1] "Column number: 16. Variable name: BDIY_log"
## [1] "Column number: 17. Variable name: SPX_log"
## [1] "Column number: 18. Variable name: DX1_log"
## [1] "Column number: 19. Variable name: GC1_log"
## [1] "Column number: 20. Variable name: HO1_log"
## [1] "Column number: 21. Variable name: USCI_log"
## [1] "Column number: 22. Variable name: GNR_log"
## [1] "Column number: 23. Variable name: SHCOMP_log"
## [1] "Column number: 24. Variable name: FTSE_log"
## [1] "Column number: 25. Variable name: respLAG"

```

---

2) Develop a process that allows us to add additional drivers which are transformations of the raw input timeseries. Include the following transformations:

- a. Rolling standard deviation (of arbitrary window)
- b. Rolling mean (of arbitrary window)
- c. Lagging (of arbitrary order)

- d. Leading (of arbitrary order)
- e. Differencing
- f. Spread (between two input drivers)
- g. Ratio (between two input drivers)
- h. Product (between two input drivers)

- Function `data_trans` includes:

`roll_std_dev` (Rolling standard deviation) - 7 days selected `roll_mean` (Rolling mean) - 7 days selected `lag_1` (Lagging) - order 1 selected `lead` (Leading) - order 1 selected `diff` (Differencing)

```
data_trans <- function(raw_data_1) {

  # add input driver to dataframe
  df_1 <- as.data.frame(raw_data_1)
  oil_data_1 <- raw_data_1
  oil_data_1 <- as.matrix(oil_data_1)

  # Rolling standard deviation, window = 7
  roll_std_dev <- roll::roll_sd(oil_data_1, 7)
  df_1$roll_std_dev <- roll_std_dev

  # Rolling mean, window = 7
  roll_mean <- roll::roll_mean(oil_data_1, 7)
  df_1$roll_mean <- roll_mean

  # Lagging, order = 1
  df_1$lag_1 <- dplyr::lag(raw_data_1)

  # Leading, order = 1
  df_1$lead <- dplyr::lead(raw_data_1)

  # Differencing
  Diff <- raw_data_1 %>% diff()
  Diff[1000] <- NA
  df_1$diff <- Diff

  return(df_1)
}

oil_data_trans <- data_trans(oil_data$OILPRICE)
head(oil_data_trans, n = 10)
```

##	raw_data_1	roll_std_dev	roll_mean	lag_1	lead	diff
## 1	4.640923	NA	NA	NA	4.633077	-0.0078462165
## 2	4.633077	NA	NA	4.640923	4.634049	0.0009720063
## 3	4.634049	NA	NA	4.633077	4.646312	0.0122629838
## 4	4.646312	NA	NA	4.634049	4.631520	-0.0147921680
## 5	4.631520	NA	NA	4.646312	4.627616	-0.0039035865
## 6	4.627616	NA	NA	4.631520	4.635214	0.0075979325
## 7	4.635214	0.006223043	4.635530	4.627616	4.635796	0.0005820722
## 8	4.635796	0.005767563	4.634798	4.635214	4.640055	0.0042582083
## 9	4.640055	0.006018100	4.635795	4.635796	4.645544	0.0054894923
## 10	4.645544	0.006957400	4.637437	4.640055	4.649665	0.0041213457

You can pass any dataframe column to the function. It will calculate all the compositions

- Building another function called `data_trans_2` to deal with 2 drivers. It contains the Ratio and Product. I couldn't solve the spread between two input drivers. Maybe I needed more time to invest in a satisfactory answer.

- f. Spread (between two input drivers)
- g. Ratio (between two input drivers)
- h. Product (between two input drivers)

```
data_trans_2 <- function(raw_data_2){  
  
  # add input driver to dataframe  
  df_2 <- as.data.frame(raw_data_2)  
  
  # Ratio bet  
  df_2$Ratio <- df_2[,1]/df_2[,2]  
  
  # Product  
  df_2$Product <- df_2[,1] * df_2[,2]  
  
  return(df_2)  
}  
  
df_2 <- data_trans_2(oil_data[, c("OILPRICE", "respLAG")])  
head(df_2, n = 10)
```

```
##      OILPRICE respLAG      Ratio Product  
## 1  4.640923 4.631812 1.0019671 21.49589  
## 2  4.633077 4.640923 0.9983093 21.50176  
## 3  4.634049 4.633077 1.0002098 21.46991  
## 4  4.646312 4.634049 1.0026463 21.53124  
## 5  4.631520 4.646312 0.9968164 21.51949  
## 6  4.627616 4.631520 0.9991572 21.43290  
## 7  4.635214 4.627616 1.0016419 21.44999  
## 8  4.635796 4.635214 1.0001256 21.48791  
## 9  4.640055 4.635796 1.0009185 21.51035  
## 10 4.645544 4.640055 1.0011831 21.55558
```

You can pass any dataframe column to the function. It will calculate the ratio and the product between the two columns

- 3) We must be able to have composition of transformations. Example: First calculate the difference between OILPRICE and resp\_LAG, and then calculate the rolling standard deviation.

- Building `data_composition` function. This also works with two columns.

```
data_composition <- function(raw_data_3){  
  
  # add input driver to dataframe  
  df_3 <- as.data.frame(raw_data_3)
```

```

# Difference
difference <- (df_3$OILPRICE - df_3$respLAG)
df_3$difference <- difference
difference <- as.matrix(difference)

# Rolling standard deviation, window = 7
roll_std <- roll::roll_sd(difference, 7)
df_3$roll_std <- roll_std

return(df_3)
}

df_3 <- data_composition(oil_data[,c("OILPRICE", "respLAG")])

head(df_3, n = 10)

```

```

##      OILPRICE respLAG      difference      roll_std
## 1  4.640923 4.631812  0.0091112388           NA
## 2  4.633077 4.640923 -0.0078462165           NA
## 3  4.634049 4.633077  0.0009720063           NA
## 4  4.646312 4.634049  0.0122629838           NA
## 5  4.631520 4.646312 -0.0147921680           NA
## 6  4.627616 4.631520 -0.0039035865           NA
## 7  4.635214 4.627616  0.0075979325 0.009882852
## 8  4.635796 4.635214  0.0005820722 0.009140087
## 9  4.640055 4.635796  0.0042582083 0.008704563
## 10 4.645544 4.640055  0.0054894923 0.008868343

```

- You can pass any dataframe column to the function. It will calculate the ratio and the product between the two columns

---

4) The sequence of transformations, and which drivers they act on must be specified by the user. One of the main purposes of this challenge is to develop a generic framework to allow this.

- While calling the previously created functions, the user need to select the correct input drivers. Then select the sequence of transformation on the `final_drivers` variable below.

```

final_drivers <- cbind(oil_data_trans, df_2, df_3)
final_drivers <-
  final_drivers[, c("raw_data_1", "roll_std_dev", "roll_mean",
                    "lag_1", "lead", "diff", "Ratio", "Product",
                    "roll_std")]

head(final_drivers, n = 10)

```

```

##      raw_data_1 roll_std_dev roll_mean      lag_1      lead      diff      Ratio
## 1    4.640923          NA          NA      NA 4.633077 -0.0078462165 1.0019671
## 2    4.633077          NA          NA 4.640923 4.634049  0.0009720063 0.9983093
## 3    4.634049          NA          NA 4.633077 4.646312  0.0122629838 1.0002098

```

```
## 4      4.646312      NA      NA 4.634049 4.631520 -0.0147921680 1.0026463
## 5      4.631520      NA      NA 4.646312 4.627616 -0.0039035865 0.9968164
## 6      4.627616      NA      NA 4.631520 4.635214 0.0075979325 0.9991572
## 7      4.635214 0.006223043 4.635530 4.627616 4.635796 0.0005820722 1.0016419
## 8      4.635796 0.005767563 4.634798 4.635214 4.640055 0.0042582083 1.0001256
## 9      4.640055 0.006018100 4.635795 4.635796 4.645544 0.0054894923 1.0009185
## 10     4.645544 0.006957400 4.637437 4.640055 4.649665 0.0041213457 1.0011831
##      Product      roll_std
## 1  21.49589      NA
## 2  21.50176      NA
## 3  21.46991      NA
## 4  21.53124      NA
## 5  21.51949      NA
## 6  21.43290      NA
## 7  21.44999 0.009882852
## 8  21.48791 0.009140087
## 9  21.51035 0.008704563
## 10 21.55558 0.008868343
```

---

- 5) For all drivers, either in their raw form or those that results from the application of one or several transformations, we must keep a meta data object where the sequence of transformations is stored. This will allow us to keep track of the meaning of each new driver. Combine all the drivers from their raw form or those that result from the application of one or several transformations using `cbind()`, named dataset as `final_drivers`.

- Creating meta data object:

```
# labeling the variables
print_with_label <- function(dframe){
  stopifnot(inherits(dframe, "data.frame"))
  labs <- labelVector::get_label(dframe, names(dframe))
  labs <- sprintf("%s: %s", names(dframe), labs)
  #print(dframe)
  cat("\n")
  cat(labs, sep = "\n")
}
final_drivers <- set_label(final_drivers,
  raw_data_1 = "target variable",
  roll_std_dev = "Rolling standard deviation(window = 7)",
  roll_std = "Rolling standard deviation(window = 7)",
  roll_mean = "Rolling mean (window = 7)",
  lag_1 = "Lagging (order = 1)",
  lead = "Leading (order = 1)",
  diff = "Differencing (order = 1)",
  Ratio = "Ration between two input drivers",
  Product = "Multiplication between two input drivers"
)
```

```
contents(final_drivers)
```

```
##
```

```
## Data frame:final_drivers 1000 observations and 9 variables      Maximum # NAs:6
##
##
##
##              Labels      Class Storage NAs
## raw_data_1          target variable numeric   double    0
## roll_std_dev    Rolling standard deviation(window = 7)  matrix   double    6
## roll_mean              Rolling mean (window = 7)  matrix   double    6
## lag_1              Lagging (order = 1) numeric   double    1
## lead              Leading (order = 1) numeric   double    1
## diff              Differencing (order = 1) numeric   double    1
## Ratio              Ration between two input drivers numeric   double    0
## Product            Multiplication between two input drivers numeric   double    0
## roll_std          Rolling standard deviation(window = 7)  matrix   double    6
```

```
k <- contents(final_drivers, sort='names', prlevels=FALSE)
print(k)
```

```
##
## Data frame:final_drivers 1000 observations and 9 variables      Maximum # NAs:6
##
##
##
##              Labels      Class Storage NAs
## raw_data_1          target variable numeric   double    0
## roll_std_dev    Rolling standard deviation(window = 7)  matrix   double    6
## roll_mean              Rolling mean (window = 7)  matrix   double    6
## lag_1              Lagging (order = 1) numeric   double    1
## lead              Leading (order = 1) numeric   double    1
## diff              Differencing (order = 1) numeric   double    1
## Ratio              Ration between two input drivers numeric   double    0
## Product            Multiplication between two input drivers numeric   double    0
## roll_std          Rolling standard deviation(window = 7)  matrix   double    6
```

```
# saving metadata.csv
lapply(k, function(x) write.table( data.frame(x), 'metadata.csv' , append= T, sep=', ' ))
```

```
## $contents
## NULL
##
## $dim
## NULL
##
## $maxnas
## NULL
##
## $id
## NULL
##
## $rangevar
## NULL
##
## $valuesvar
## NULL
##
```

```
## $unique.ids
## NULL
##
## $range
## NULL
##
## $values
## NULL
##
## $dfname
## NULL
##
## $Levels
## NULL
##
## $longLabels
## NULL
```

---

6) For each driver that results from the user-specified sequence of transformations, we need to assess a few statistics: Normality test Stationarity test Correlation coefficient with the target These statistics need to be stored in the meta data object. The purpose of this is, we may be interested in keeping in the final model only drivers that are normally distributed, or only drivers whose correlation with the target is above a given threshold, or another combination of such criteria.

- Normality test.

```
# normality graph
normality <- function(input_driver, p_value) {

  print(ggdensity(input_driver,
    main = "Density plot of Rolling Standard deviation",
    xlab = "",
    add = 'mean',
    ggtheme = theme_classic(),
    rug = TRUE))

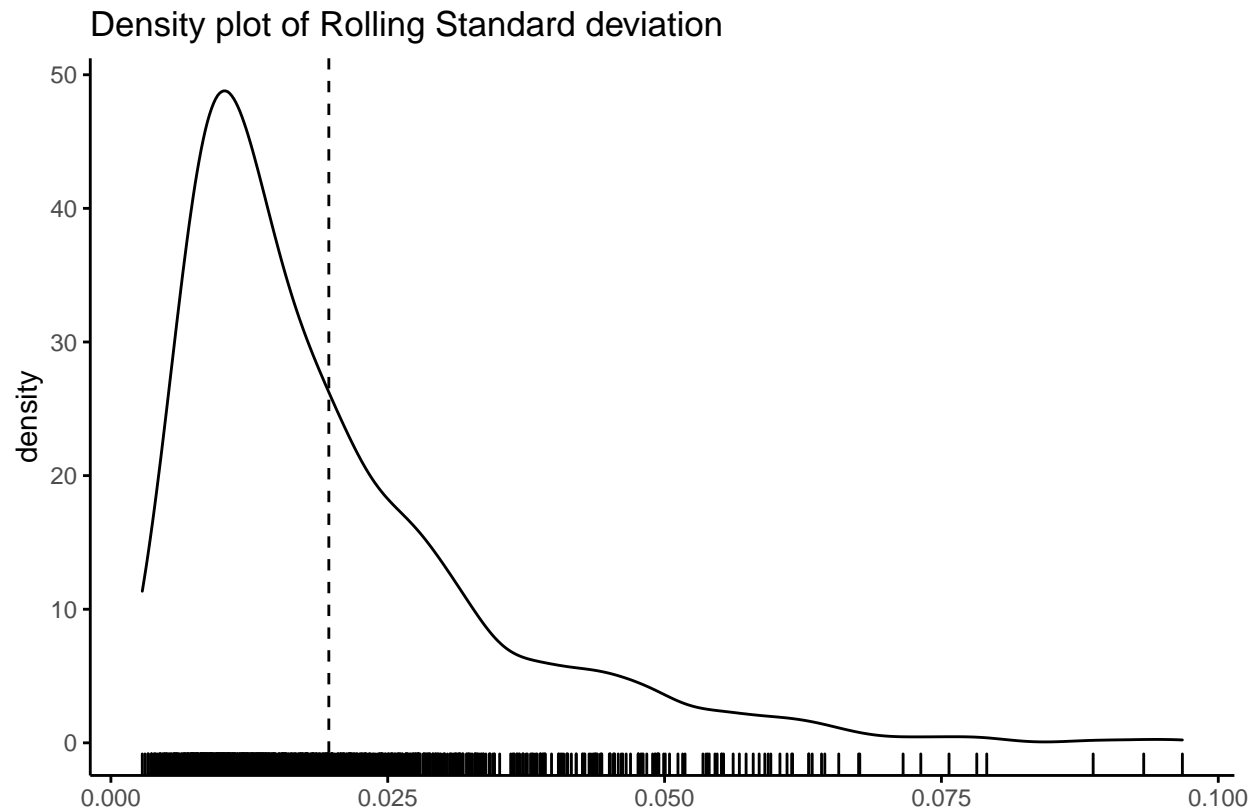
  z <- shapiro.test(input_driver)
  print(shapiro.test(input_driver))

  if(z[2]>= p_value){
    print('Normally Distributed')
    x <- sys.call()
    x <- as.character(x)
    norm_lst <- append(norm_lst, x)
  }
  else{
    print('Not normally distributed')
  }
}

# function call
normality(final_drivers$roll_std_dev, 0.05)
```



```
## Don't know how to automatically pick scale for object of type labelled/matrix/array. Defaulting to c
```



```
##  
## Shapiro-Wilk normality test  
##  
## data: input_driver  
## W = 0.84009, p-value < 2.2e-16  
##  
## [1] "Not normally distributed"
```

If the result of the p-value is higher or equal to the passed p-value, the name of the variable is saved on `norm_lst` and stored in a meta data object. Usually,  $p\text{-value} \leq 0.05$  means that the distribution is significantly different than normal distribution.

---

#### b. Stationarity test

- Augmented Dickey-Fuller (ADF) t-statistic is used to find if the series has a unit root (a series with a trend line will have a unit root and result in a large p-value). If the  $p\text{-value} < 0.05$  then data is stationary if  $p\text{-value} > 0.05$  then data is non-stationary.

Before the test, we remove NA values and replace them with 0.

```

# Stationarity check

stationarity <- function(input_driver, p_value) {
  input_driver[is.na(input_driver)] <- 0
  tseries::adf.test(input_driver)
  sz <- tseries::adf.test(input_driver)

  if(sz[2]<= p_value){
    y <- sys.call()
    y <- as.character(y)
    stat_lst <- append(norm_lst, y)
    print(sz)
    print('Stationary Data')
  }
  else{
    print(sz)
    print('Non-stationary Data')
  }
}

# function call
stationarity(final_drivers$roll_mean, 0.05)

```

```

##
## Augmented Dickey-Fuller Test
##
## data: input_driver
## Dickey-Fuller = -1.7106, Lag order = 9, p-value = 0.7008
## alternative hypothesis: stationary
##
## [1] "Non-stationary Data"

```

- To test another drivers, just replace the input\_driver. The data is also stored in a metadata object

```

# saving normal data in metadata_normality.csv
try(lapply(norm_lst, function(x) write.table( data.frame(x), 'metadata_normality.csv' , append= T, sep=',')))

# saving stationary data in stationary_normality.csv
try(lapply(stat_lst, function(x) write.table( data.frame(x), 'metadata_stationary.csv' , append= T, sep=',')))

```

---

c. Correlation coefficient with the target

```

# Correlation coefficient
correlation <- function(input_drivers){

  input_drivers[is.na(input_drivers)] <- 0
  corr_mat=cor(input_drivers)
  col <- colorRampPalette(c("#BB4444", "#EE9988", "#FFFFFF", "#77AADD", "#4477AA"))
  return(corrplot(corr_mat, method="color",

```

```

    type="upper", order="hclust",
    addCoef.col = "black",
    tl.col="black", tl.srt=45,
    # hide correlation coefficient on the principal diagonal
    diag=FALSE
  ))
}

correlation(final_drivers)

```

